



SPECULATIVE BARRIERS WITH TRANSACTIONAL MEMORY

M. PEDRERO, R. QUISLANT, E. GUTIÉRREZ, E. L. ZAPATA, O. PLATA

DEPT. COMPUTER ARCHITECTURE




UNIVERSIDAD
DE MÁLAGA

CONTENTS


- **Background**
- Introduction
- Speculative Barriers (SBs)
 - API
 - SBs and non-transactional codes
- Experimental Evaluation
 - Methodology
 - Results
- Related Work
- Conclusions

BACKGROUND

LOCKS




TRANSACTIONS




BACKGROUND

LOCKS




TRANSACTIONS




BACKGROUND

LOCKS



TRANSACTIONS



BACKGROUND

■ Hardware Transactional Memory [Herlihy'93]

■ ISA extensions:

- Load (exclusive) transactional, store transactional, abort, validate, commit
- Explicit HTM: each load/store is to be annotated

■ Per-core small private transactional fully-associative cache:

- Aborts/commits in one-cycle
- Small/short transactions

■ Coherence protocol modification:

- In charge of detecting conflicts among transactions

■ Pros & Cons:

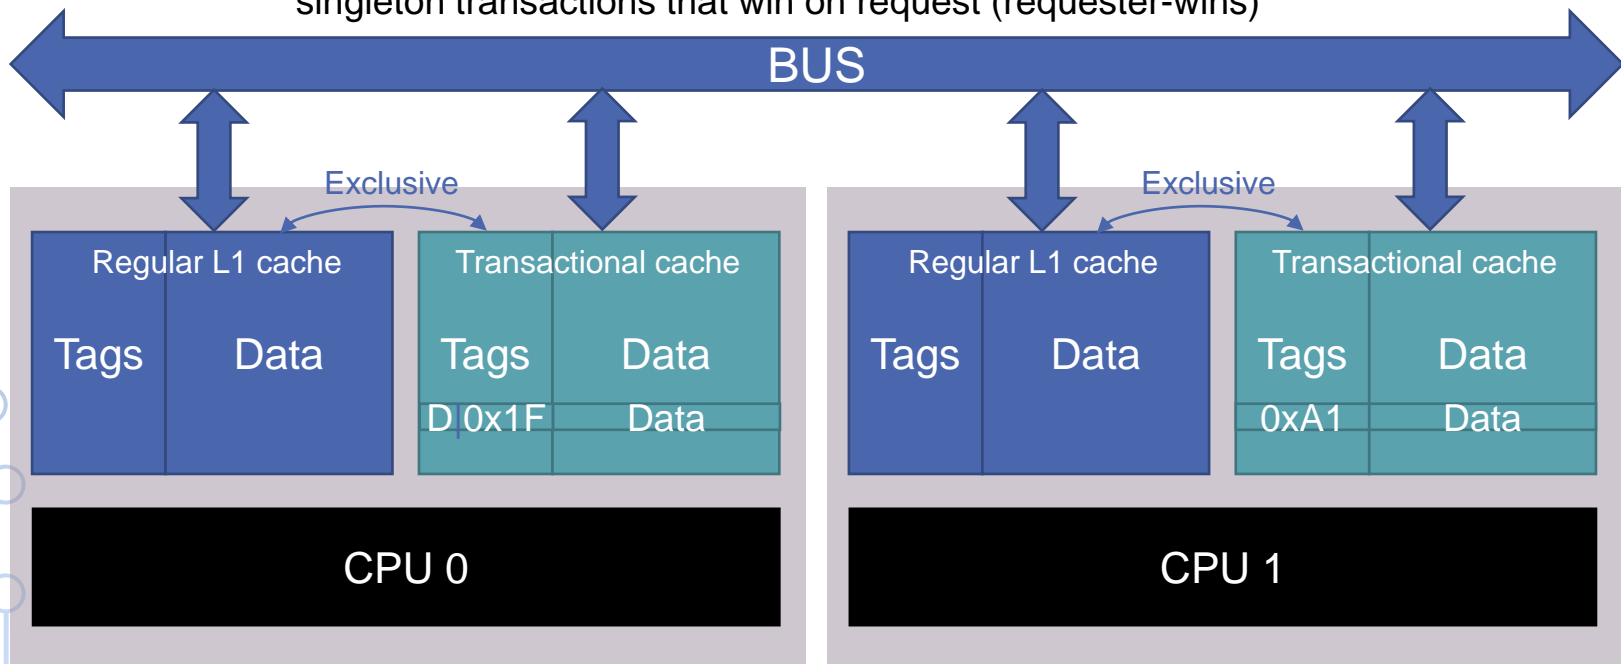
- Performance: fewer accesses
- Forward progress not guaranteed

```
ST(&counter, LTX(&counter) + 1);  
if (COMMIT()) { Transaction  
    success++;  
    backoff = BACKOFF_MIN;  
} else {  
    wait = random() % (1<<backoff);  
    while (wait--);  
    if (backoff < BACKOFF_MAX)  
        backoff++;  
}
```

BACKGROUND

■ Hardware Transactional Memory [Herlihy'93]

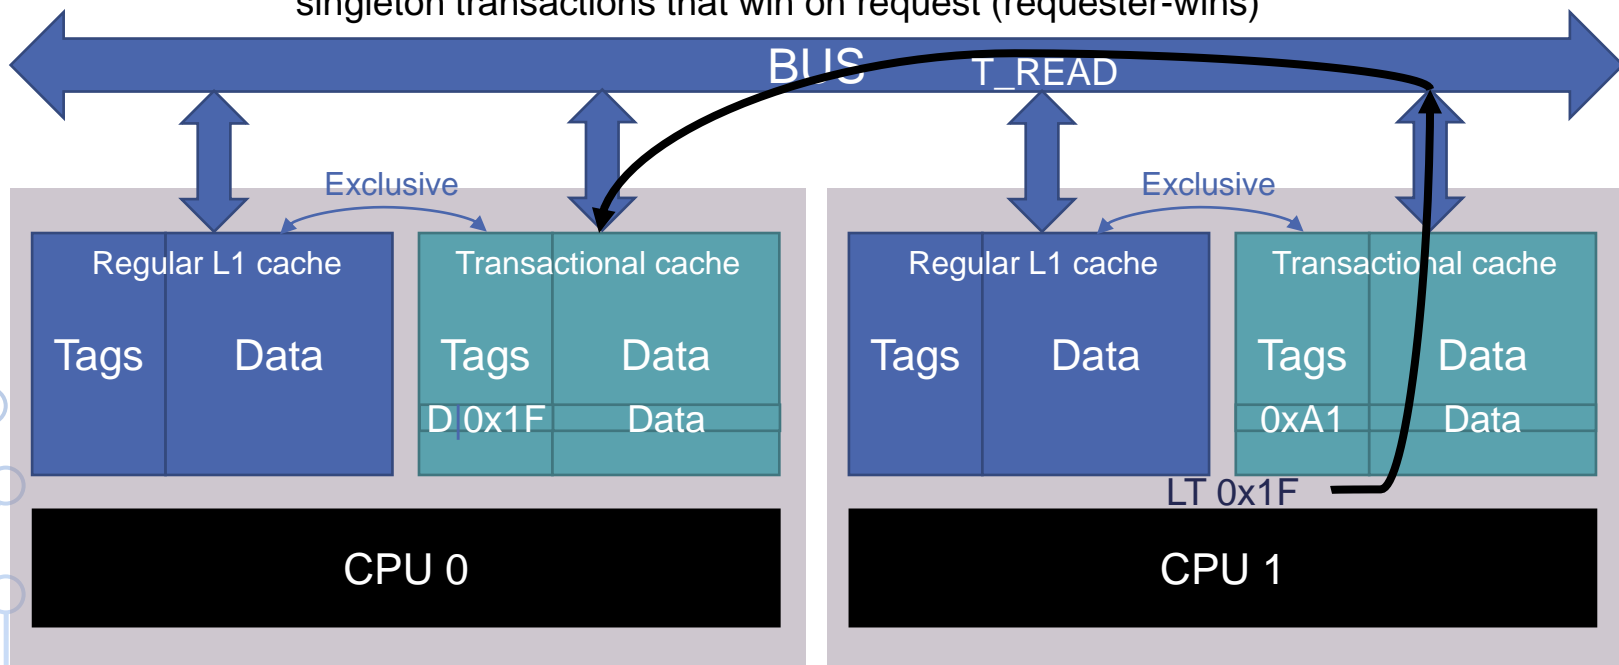
- Conflict resolution policy: Requester aborts
- Herlihy's approach does not define behavior on interaction with non-transactional code
 - Strong isolation [Blundell'06]: non-transactional instructions are treated as singleton transactions that win on request (requester-wins)



BACKGROUND

■ Hardware Transactional Memory [Herlihy'93]

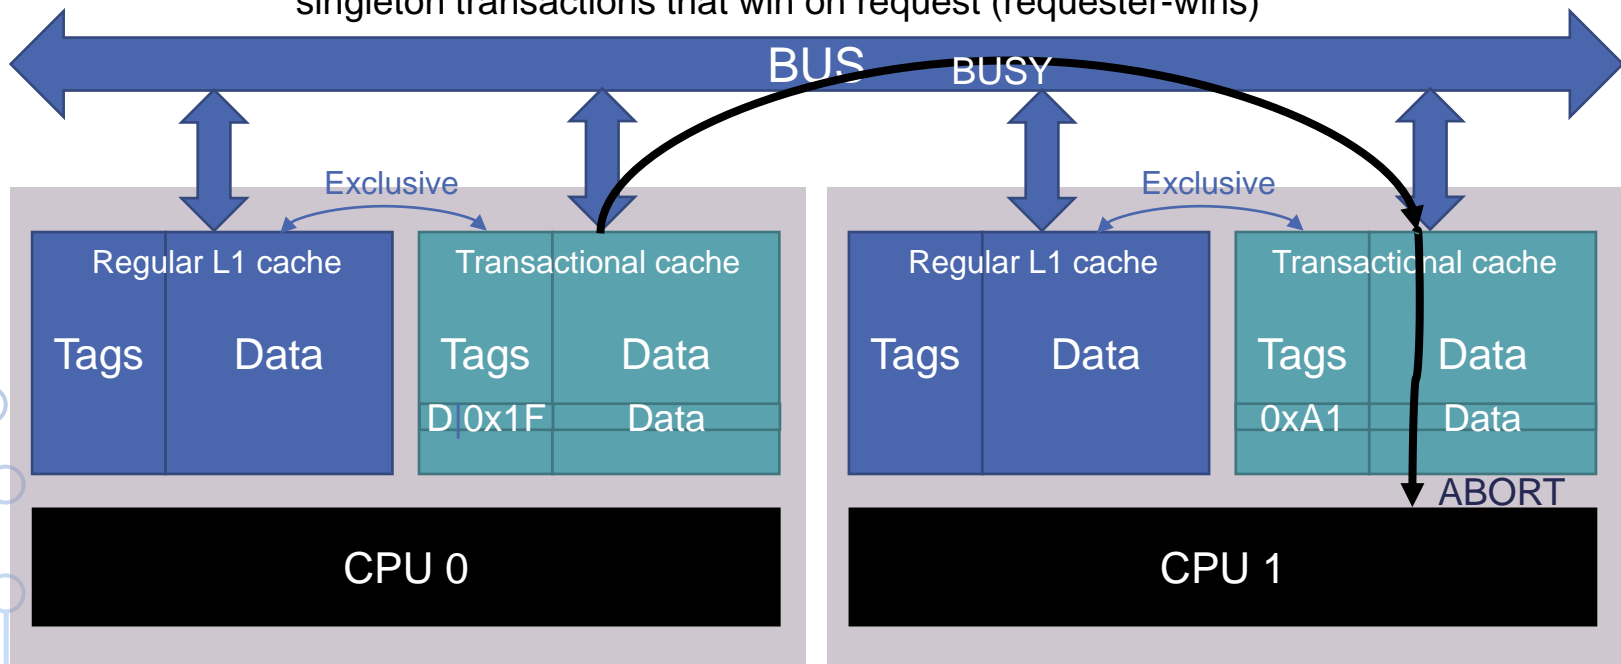
- Conflict resolution policy: Requester aborts
- Herlihy's approach does not define behavior on interaction with non-transactional code
 - Strong isolation [Blundell'06]: non-transactional instructions are treated as singleton transactions that win on request (requester-wins)



BACKGROUND

■ Hardware Transactional Memory [Herlihy'93]

- Conflict resolution policy: Requester aborts
- Herlihy's approach does not define behavior on interaction with non-transactional code
 - Strong isolation [Blundell'06]: non-transactional instructions are treated as singleton transactions that win on request (requester-wins)



BACKGROUND

■ Research following Herlihy's seminal paper:

- Focused on supporting larger/longer transactions:
 - Transactional data set (read/write set) is limited to cache size
 - Transactions do not survive interrupts
- Proposals:
 - UTM [Ananian'05]: Transactions nearly as large as virtual memory; transaction information in memory log (surviving context switches and migrations); nested transactions are flattened
 - VTM [Rajwar'05]: Overflowed transactions allocate entries in a virtual memory data structure; page faults are managed by suspending-resuming the transaction
 - LogTM-SE [Yen'07]: Software undo-log; signatures for fast/unbound conflict detection; escape actions; closed-nesting with partial aborts and open-nesting
 - OneTM [Blundell'07]: Allows for one overflowed transaction to continue, while others stall

BACKGROUND

■ Commercial HTM extensions:

- Back to Herlihy's **best-effort** approach
- **Intel TSX** (Transactional Synchronization eXtensions), Core Haswell, 2013:
 - HLE: Hardware Lock Elision
 - Prefixes to certain instructions (e.g., XCHG) to elide locks and execute transactions instead
 - Enables backwards compatibility and portability
 - RTM: Restricted Transactional Memory
 - New instructions (`xbegin`, `xend`, `xabort`): Implicit HTM
 - L1D cache holds transactional stores (~22KB); L3 keeps track of transactional loads (~4MB) [Nakaike'15]
 - Cache coherence protocol implements a requester-wins conflict resolution policy
 - Flat nesting
 - Transaction **suspend/resume** (escape actions) is **not supported**
 - Future extensions (Sapphire Rapids) to include Suspend Load Address Tracking (TSXLDTRK)

BACKGROUND

■ Commercial HTM extensions:

- **Intel TSX Transaction API:** wrappers implementing retry and fallback path

```
#define XBEGIN()  
    int retries = 0;  
    do {  
        retries++;  
        if (retries > MAX_RETRIES) {  
            acquire_lock(g_fallback_lock);  
            break;  
        }  
    } while (_xbegin() != _XBEGIN_STARTED)
```

```
#define XEND()  
    if (retries <= MAX_RETRIES) {  
        if (g_fallback_lock) _xabort(LOCK_TAKEN);  
        _xend();  
    } else  
        release_lock(g_fallback_lock);
```

```
__inline unsigned int _xbegin() {  
    unsigned status;  
    __asm {  
        move eax, _XBEGIN_STARTED  
        xbegin _txnL1  
_txnL1: move status, eax  
    }  
    return status;  
}
```

Intrinsic

BACKGROUND

■ Commercial HTM extensions:

■ IBM Power8/9, 2014/2017:

- New instructions (`tbegin`, `tend`, `tabort`, `tsuspend`, `tresume`): Implicit HTM
- L2 TMCAM (Content Addressable Memory) holds transactional loads and stores (8KB) [Nakaike'15]
- Flat nesting
- Transaction **suspend/resume** (escape actions) **are supported**
 - Debugging
 - Thread-Level Speculation: extracting parallelism from a sequential application
 - One thread expanded per loop iteration
 - Loop iteration enclosed within a transaction to detect inter-loop dependencies
 - Escaped loop on a control variable before calling `tend`
 - **Speculative Barriers**

■ ARM TME (Transactional Memory Extension), ARMv9, 2021

BACKGROUND

■ Commercial HTM extensions:

- **Power8 Transaction API:** wrappers implementing retry and fallback path

```
#define TBEGIN()  
    int retries = 0;  
    do {  
        retries++;  
        if (retries > MAX_RETRIES) {  
            acquire_lock(g_fallback_lock);  
            break;  
        }  
    } while (!_tbegin())
```

Returns true if the transaction was successfully started; and false if not, or on abort

```
#define TEND()  
    if (retries <= MAX_RETRIES) {  
        if (g_fallback_lock) _xabort(LOCK_TAKEN);  
        _tend();  
    } else  
        release_lock(g_fallback_lock);
```

CONTENTS

- Background
- **Introduction**
- Speculative Barriers (SBs)
 - API
 - SBs and non-transactional codes
- Experimental Evaluation
 - Methodology
 - Results
- Related Work
- Conclusions

INTRODUCTION


■ Speculative Barriers (SB)

- A thread reaching a barrier elides it and executes transactionally from there on
- The transaction opened by a thread after crossing a barrier (SB transaction) is committed as soon as all threads cross the barrier
- Escape actions are used inside transactions to perform control actions and prevent aborts
- A partial order is implemented among transactions before and after a barrier
- We provide an API for SBs including wrappers for beginning/committing a transaction, and a barrier substitute
 - IBM Power8 HTM extensions are used for the implementation (tsuspend-tresume)

INTRODUCTION

■ Motivation

- *Target:* barrier-intensive applications that exhibits poor performance due to:
 - Load imbalance (idle threads)
 - Contention (many threads, one barrier lock)
- *Solution:* Harnessing imbalance to get ahead with work and eliding the barrier to avoid contention



Standard barriers (top) versus SBs (bottom).


INTRODUCTION

■ Motivation: Recurrence

```

chunk = N/#TH;
start = tid*chunk;
stop = MIN(start+chunk,N);
for (t=0; t ≤ N-2; t++) {
    for (k=start; k<stop; k++)
    {
        if (k < (n-t-1))
            w[t+k+1] +=
                b[k][t+k+1]*w[t];
    }
    barrier();
}

```



CONTENTS

- Background
- Introduction
- **Speculative Barriers (SBs)**
 - **API**
 - **SBs and non-transactional codes**
- Experimental Evaluation
 - Methodology
 - Results
- Related Work
- Conclusions

SPECULATIVE BARRIERS

■ SB implementation:


- SBs first intended for codes with transactions
- We use begin/commit as well as barrier wrappers to implement SBs

```

chunk = N/#TH;
start = tid*chunk;
stop = MIN(start+chunk,N);
for (t=0; t ≤ N-2; t++) {
    for (k=start; k<stop; k++) {
        TX_START();
        if (k < (n-t-1))
            w[t+k+1] +=
                b[k][t+k+1]*w[t];
        TX_STOP();
    }
    SPEC_BARRIER();
}
LAST_BARRIER();

```

Recurrence



SPECULATIVE BARRIERS

- **SB implementation: SB API**
 - Global variables and transaction descriptor

```
global int sbBarrier = #TH; //Barrier counter
global int sbOrder = 1; //Barrier global order
thread local tx {
    int order = 1; //Local order
    int retries = 0;
    bool spec = 0; //Speculation flag
    int specMax = MAX_SPEC; //Adaptive
    int specLevel = MAX_SPEC; //speculation level
}
```

SPECULATIVE BARRIERS

■ SB implementation: SB API

■ TX_START() wrapper

```
chunk = N/#TH;
start = tid*chunk;
stop = MIN(start+chunk,N);
for (t=0; t ≤ N-2; t++) {
    for (k=start; k<stop; k++) {
        TX_START();
        if (k < (n-t-1))
            w[t+k+1]+=
                b[k][t+k+1]*w[t];
        TX_STOP();
    }
    SPEC_BARRIER();
}
LAST_BARRIER();
```

Recurrence

```
#define TX_START()
if (tx.spec == 0) { //Not executing an SB
    TBEGIN(); //Normal transaction begin
} //If in SB mode => do nothing
```

Transaction API

```
#define TBEGIN()
int retries = 0;
do {
    retries++; → tx.retries
    if (retries > MAX_RETRIES) {
        acquire_lock(g_fallback_lock);
        break;
    }
} while (!_tbegin())
```

SPECULATIVE BARRIERS

■ SB implementation: SB API

■ TX_STOP() wrapper

```

chunk = N/#TH;
start = tid*chunk;
stop = MIN(start+chunk,N);
for (t=0; t ≤ N-2; t++) {
    for (k=start; k<stop; k++) {
        TX_START();
        if (k < (n-t-1))
            w[t+k+1]+=
                b[k][t+k+1]*w[t];
        TX_STOP();
    }
    SPEC_BARRIER();
}
LAST_BARRIER();

```

Recurrence

```

#define TX_STOP()
if (tx.spec == 0) { //Not in SB mode
    TEND(); //Normal transaction end
    resetDescriptor(tx);
} else { //In SB mode
    _tsuspend();
    if (tx.order ≤ sbOrder) {
        //1. All crossed the barrier
        _tresume();
        _tend(); //Commit SB
        resetDescriptor(tx);
    } else {
        //2. Some threads didn't cross
        _tresume();
        tx.specLevel--;
        if (tx.specLevel == 0) {
            _tsuspend();
            while (tx.order > sbOrder) ;
            _tresume();
            _tend(); //Commit SB
            resetDescriptor(tx);
        }
    }
}

```

Busy wait

SPECULATIVE BARRIERS

■ SB implementation: SB API

■ SPEC_BARRIER() wrapper

```

chunk = N/#TH;
start = tid*chunk;
stop = MIN(start+chunk,N);
for (t=0; t ≤ N-2; t++) {
    for (k=start; k<stop; k++) {
        TX_START();
        if (k < (n-t-1))
            w[t+k+1]+=
                b[k][t+k+1]*w[t];
        TX_STOP();
    }
    SPEC_BARRIER();
}
LAST_BARRIER();

```

Recurrence

```

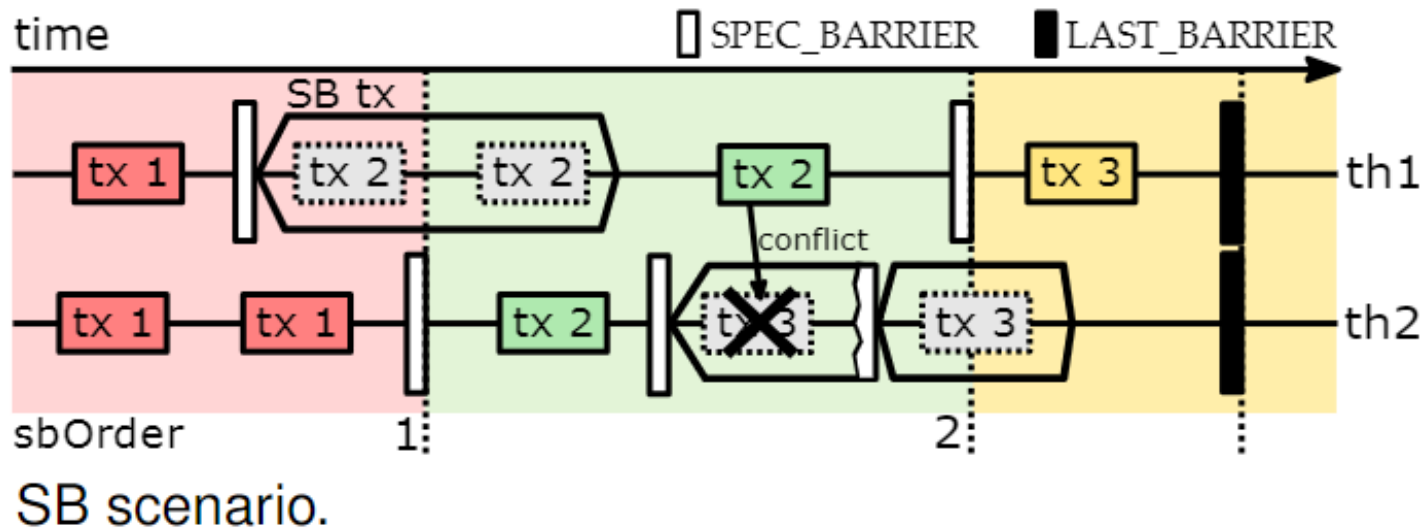
#define SPEC_BARRIER()
if (tx.spec == 1) {
    _tsuspend();
    while (tx.order > sbOrder) ;
    _tresume();
    _tend(); //Commit SB
    resetDescriptor(tx);
}
tx.order++; //Local order
if (_addAndFetch(sbBarrier,-1) == 0) {
    sbBarrier = #TH; //Barrier reset
    _addAndFetch(sbOrder,1); //Global order
} else { //Not the last thread to cross
    do {
        tx.retries++;
        if (tx.order ≤ sbOrder) {
            resetDescriptor(tx);
            break;
        }
        tx.spec = 1;
        decrementMaxSpecLevel();
    } while (!_tbegin());
}

```


SPECULATIVE BARRIERS

■ SB implementation: SB API

- LAST_BARRIER() wrapper
 - Busy waits until speculative barrier threads
 - It is similar to SPEC_BARRIER() but it does not begin new speculative transactions



SPECULATIVE BARRIERS

■ SBs and non-transactional code

■ CHECK_SPEC() wrapper

- Transactional codes offer SB check points naturally on TX_STOP()
- In non-transactional codes, SB transactions may run for long (e.g., from barrier to barrier) → CHECK_SPEC() solves the problem


■ Implementation: like else clause from TX_STOP()

■ It can be used with transactional codes as well

- If transactions are far away each other, CHECK_SPEC() provides additional SB check points

```
for (t=0; t ≤ N-2; t++) {
  for (k=start; k<stop; k++) {
    if (k < (n-t-1))
      w[t+k+1] +=
        b[k][t+k+1]*w[t];
    CHECK_SPEC();
  }
  SPEC_BARRIER();
}
```

Recurrence



CONTENTS

- Background
- Introduction
- Speculative Barriers (SBs)
 - API
 - SBs and non-transactional codes
- **Experimental Evaluation**
 - **Methodology**
 - **Results**
- Related Work
- Conclusions

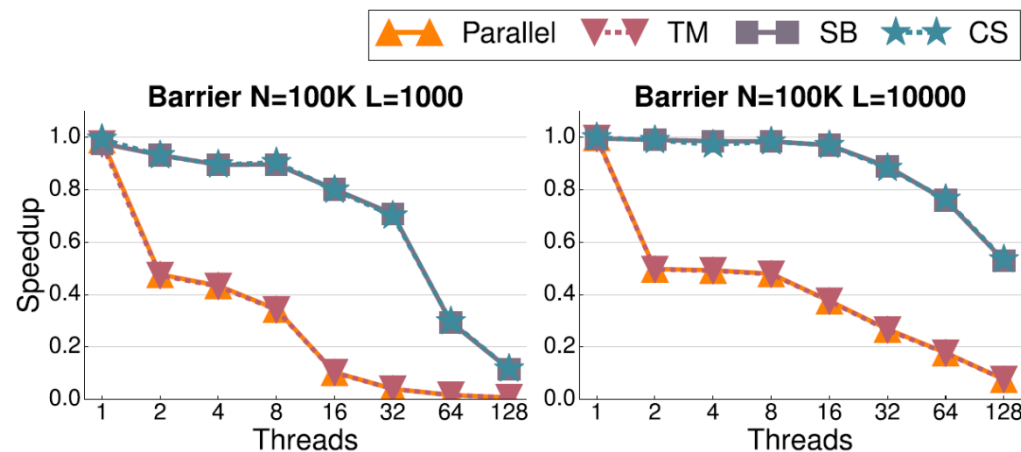
METHODOLOGY

- Server S822LC-8335 with 2 sockets:
 - IBM Power8 processor with 10 cores each and 8 SMT → 160 threads
- Best-effort HTM:
 - 8KB cache for transactional accesses
 - The cache is shared by threads when using SMT
 - 128B words
 - Escape actions (tsuspend/tresume)
- Experiments were carried out up to 128 threads with round-robin affinity:
 - 1-8 threads → 1 socket, w/o SMT
 - 16 threads → 2 sockets, w/o SMT
 - 32+ threads → 2 sockets, with SMT

METHODOLOGY

■ Experiment profiles:

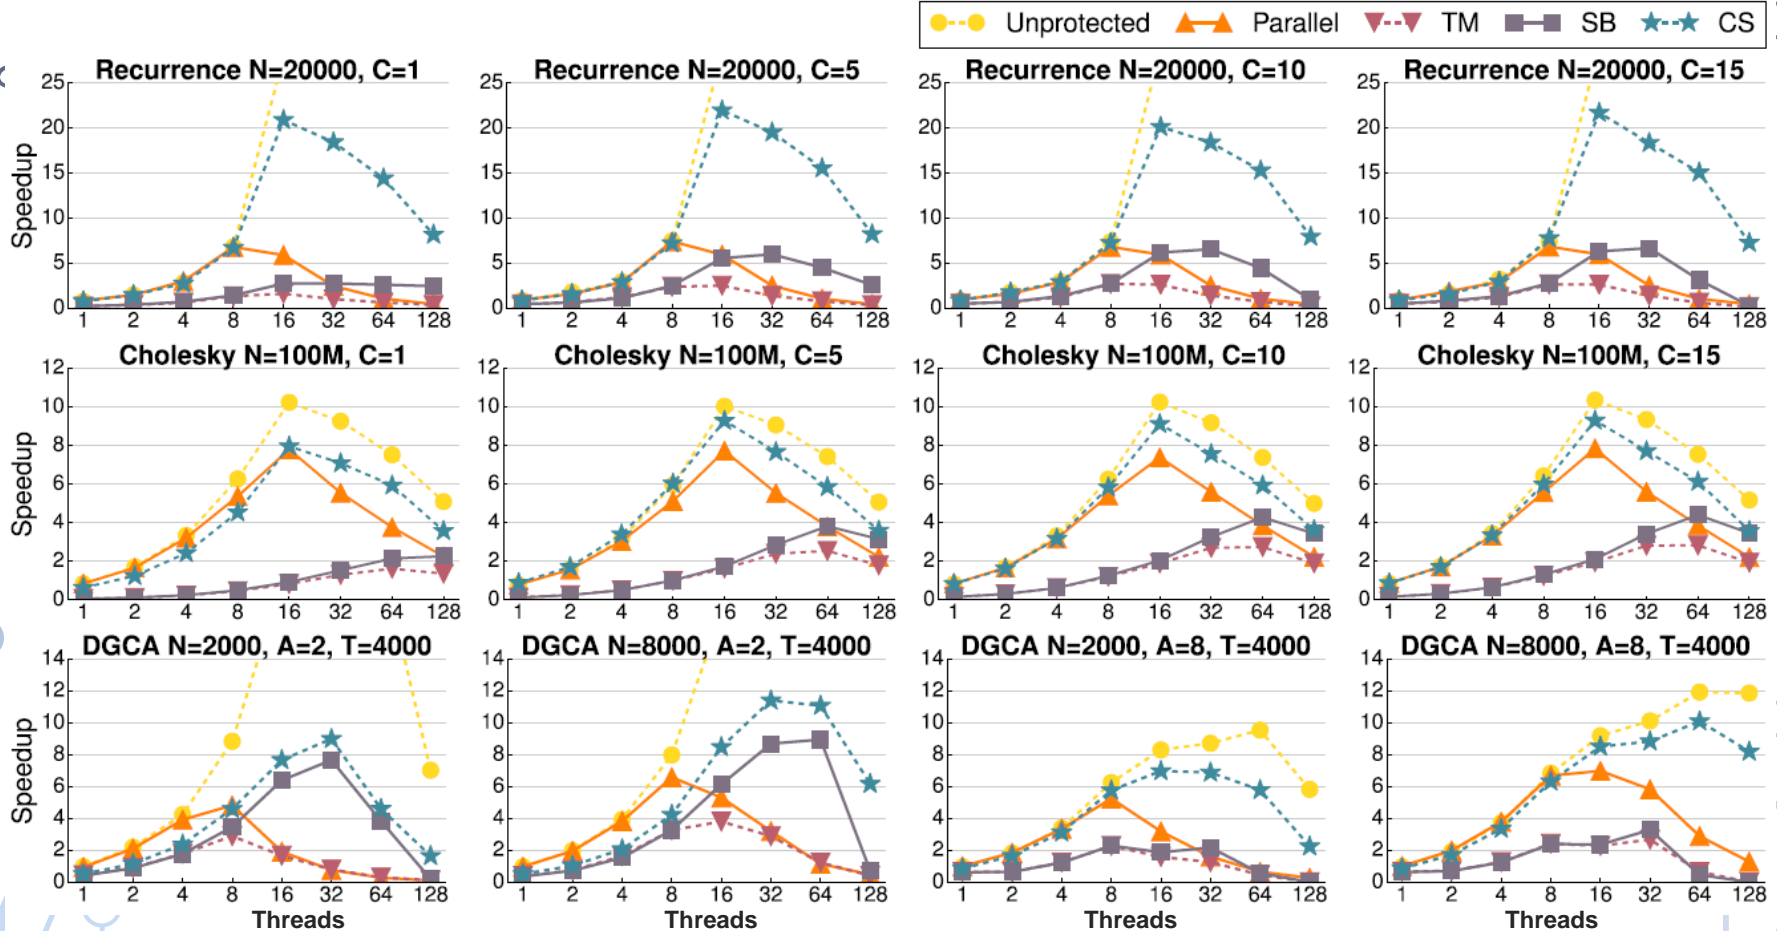
- *Unprotected*: no barriers, upper bound
- *Parallel*: no transactions and standard barriers
- *TM*: transactions and standard barriers
- *SB*: transactions and speculative barriers (SBs)
- *CS*: no transactions, SBs and CHECK_SPECS



RESULTS

	Power 8			TCR (%) / SCR (%)		
	Parameters	#Barr	tx (M/A)	TM	SB	CS
Barr	N100K L1K	100K	5/4.5	99.9/-	99.9/99.9	-/99.9
	N100K L10K	100K	5/4.5	99.9/-	99.7/99.9	-/99.9
Reccren	N20K C1	20K	9/6.5	98.6/-	99.3/6.3	-/34.6
	N20K C5	20K	16/8.9	95.5/-	98.6/8.3	-/33.1
	N20K C10	20K	21/11.8	92.2/-	97.9/8.4	-/26.7
	N20K C15	20K	26/14.8	88.6/-	96.9/7.9	-/29.2
Cholesky	N100M C1	27	8/7.2	100/-	100/38.0	-/60.9
	N100M C5	27	10/8.5	99.9/-	100/43.8	-/58.8
	N100M C10	27	12/10	99.9/-	100/45.1	-/59.3
	N100M C15	27	12/11.6	99.9/-	100/45.7	-/57.1
DGCA	N2K A2 T4K	4K	20/18.1	99.8/-	99.9/99.6	-/99
	N8K A2 T4K	4K	20/18.1	99.9/-	99.9/99.4	-/99.7
	N2K A8 T4K	4K	27/25.1	94.2/-	93.1/42.1	-/48.6
	N8K A8 T4K	4K	26/24.6	98.1/-	97.8/29.6	-/40.3
Sten	Large T100	101	12/11.1	99.9/-	99.9/10.3	-/25.5
	Large T800	801	12/11.1	99.9/-	99.9/7.6	-/26.7
Hist	Large T1K C1	2001	6/5.9	98.9/-	98.9/0.9	-/1.2
	Large T1K C5	2001	15/13.2	30.4/-	30.9/0.2	-/0.2
SSCA2	k1 -s14-l9-p9	10	8/8	98.6/-	98.5/3	-
	k1 -s20-l3-p3	10	8/8	99.9/-	99.9/0.2	-
	k4 -s14-l9-p9	24K	6/6	71.77/-	30.07/0.51	-
	k4 -s20-l3-p3	1572K	6/6	62.61/-	30.47/0.05	-

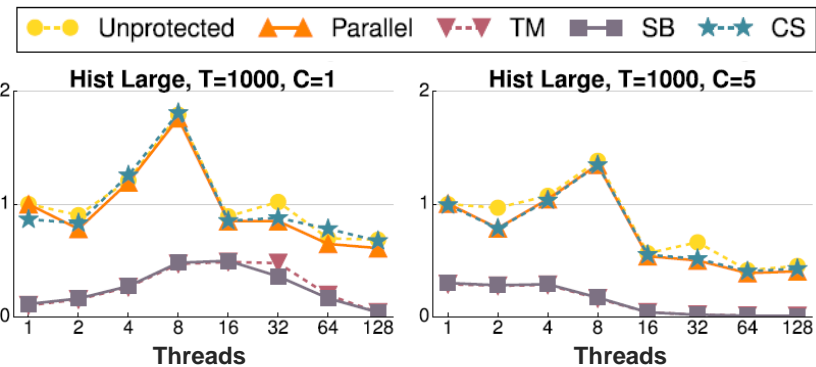

RESULTS



RESULTS



■ Bad case scenarios:

- No load imbalance
- Few barriers



■ Transactional workloads:

- Only TM and SB profiles
- More barriers and load imbalance in Kernel 4



CONTENTS

- Background
- Introduction
- Speculative Barriers (SBs)
 - API
 - SBs and non-transactional codes
- Experimental Evaluation
 - Methodology
 - Results
- **Related Work**
- Conclusions

RELATED WORK

- *[Martínez, ASPLOS'02]*
 - TLS on locks, flags and barriers
 - Dedicated hardware
 - Transaction ordering is not supported
 - Expose hardware to the user
- *[Nagarajan, LCPC'09]*
 - Two codes: with and without barriers
 - Dedicated hardware and compiler
 - 12% improvement
- *[Bonnichsen, IWOMP'15]*
 - Barrier speculation implemented with Intel's RTM
 - Transaction ordering is not supported
 - No performance gains (No escape actions)

CONTENTS

- Background
- Introduction
- Speculative Barriers (SBs)
 - API
 - SBs and non-transactional codes
- Experimental Evaluation
 - Methodology
 - Results
- Related Work
- **Conclusions**

CONCLUSIONS

- Optimistic barriers with HTM to hide latency and load imbalance
- Escape actions are used to allow communication between threads without causing aborts
- API is provided
- Improves performance in barrier-intensive applications
 - Number of threads high
 - Multi-socket configuration
- SB overload doesn't seem to penalize performance